

CS 268: Structured P2P Networks: Pastry and Tapestry

Sean Rhea
April 29, 2003

Domain

- Structured peer-to-peer overlay networks
 - Sometimes called DHTs, DOLRs, CASTs, ...
 - Examples: CAN, Chord, Pastry, Tapestry, ...
- Contrasted with unstructured P2P networks
 - Gnutella, Freenet, etc.
- Today talking about Pastry and Tapestry

2

Service Model

- Let I be a set of identifiers
 - Such as all 160-bit unsigned integers
- Let N be a set of nodes in a P2P system
 - Some subset of all possible (IP, port) tuples
- Structured P2P overlays implement a mapping
$$owner_N: I \rightarrow N$$
 - Given any identifier, deterministically map to a node
- Properties
 - Should take $O(\log |N|)$ time and state per node
 - Should be roughly load balanced

3

Service Model (con't.)

- Owner mapping exposed in variety of ways
 - In Chord, have function
$$n = find_successor(i)$$
 - In Pastry and Tapestry, have function
$$route_to_root(i, m)$$
- In general, can be iterative or recursive
 - Iterative = directed by querying node
 - Recursive = forwarded through network
- May also expose $owner^{-1}: N \rightarrow P(I)$
 - Which identifiers given node is responsible for

4

Other Service Models

- Other models can be implemented on owner
- Example: Distributed hash table (DHT)

```
void put (key, data) {  
    n = owner (key)  
    n.hash_table.insert (key, data)  
}  
  
data get (key) {  
    n = owner (key)  
    return n.hash_table.lookup (key)  
}
```

5

Lecture Overview

- Introduction
- PRR Trees
 - Overview
 - Locality Properties
- Pastry
 - Routing in Pastry
 - Joining a Pastry network
 - Leaving a Pastry network
- Tapestry
 - Routing in Tapestry
 - Object location in Tapestry
- Multicast in PRR Trees
- Conclusions

6

PRR Trees

- Work by Plaxton, Rajaraman, Richa (SPAA '97)
 - Interesting in a distributed publication system
 - Similar to Napster in interface
 - Only for static networks (set N does not change)
 - No existing implementation (AFAIK)
- Pastry and Tapestry both based on PRR trees
 - Extend to support dynamic node membership
 - Several implementations

7


PRR Trees: The Basic Idea

- Basic idea: add injective function
 $node_id: N \rightarrow I$
 - Gives each node a name in the identifier space
- $owner(i)$ = node whose $node_id$ is "closest" to i
 - Definition of closest varies, but always over I
- To find $owner(i)$ from node with identifier j
 1. Let p = longest matching prefix between i and j
 2. Find node k with longest matching prefix of $|p|+1$ digits
 3. If no such node, j is the owner (root)
 4. Otherwise, forward query to node k
- Step 2 is the tricky part

8

PRR Trees: The Routing Table

- Each node n has $O(b \log_b |N|)$ neighbors
 - Each L_x neighbor shares x digits with n
 - Set of neighbors forms a routing table



9

PRR Trees: Routing


- To find owner (47E2)
 - Query starts at node 3AF2
 - Resolve first digit by routing to 4633
 - Resolve second digit by routing to 47DA, etc.



10

PRR Trees: Routing (con't.)

- Problem: what if no exact match?
 - Consider the following network
 - Who is the owner of identifier 3701?
- Network is well formed
 - Every routing table spot that can be filled is filled
 - Can route to all node identifiers
- Owner of 3701 not well defined
 - Starting from 1000, it's node 3800
 - Starting from 2000, it's node 3600
- Violation of service model



11

PRR Trees: Handling Inexact Matches

- Want owner function to be deterministic
 - Must have a way to resolve inexact matches
- Solved different ways by each system
 - I have no idea what PRR did
 - Pastry chooses numerically closest node
 - Can break ties high or low
 - Tapestry performs "surrogate routing"
 - Chooses next highest match on per digit basis
- More on this later


12

Locality in PRR Trees

- Consider a node with id=1000 in a PRR network
 - At lowest level of routing table, node 1000 needs neighbors with prefixes 2-, 3-, 4-, etc.
 - In a large network, there may be several of each
- Idea: chose the “best” neighbor for each prefix
 - Best can mean lowest latency, highest bandwidth, etc.
- Can show that this choice gives good routes
 - For certain networks, routing path from query source to owner no more than a constant worse than routing path in underlying network
 - I’m not going to prove this today, see PRR97 for details

13

Locality in PRR Trees: Experiments



14

Lecture Overview

- Introduction
- PRR Trees
 - Overview
 - Locality Properties
- Pastry
 - Routing in Pastry
 - Joining a Pastry network
 - Leaving a Pastry network
- Tapestry
 - Routing in Tapestry
 - Object location in Tapestry
- Multicast in PRR Trees
- Conclusions

15


Pastry Introduction

- A PRR tree combined with a Chord-like ring
 - Each node has PRR-style neighbors
 - And each node knows its predecessor and successor
 - Called its leaf set
- To find owner (i), node n does the following:
 - If i is n 's leaf set, choose numerically closest node
 - Else, if appropriate PRR-style neighbor, choose that
 - Finally, choose numerically closest from leaf set
- A lot like Chord
 - Only leaf set necessary for correctness
 - PRR-neighbors like finger table, only for performance

16

Pastry Routing Example

- PRR neighbors in black
- Leaf set neighbors in blue
- Owner of 3701 is now well-defined
- From 1000
 - Resolve first digit routing to 3800
 - At 3800, see that we're done
 - (Numerically closer than 3600)
- From 2000
 - Resolve first digit routing to 3600
 - At 3600, 3701 is in leaf set
 - In range 2000-3800
 - Route to 3800 b/c numerically closer



17

Notes on Pastry Routing

- Leaf set is great for correctness
 - Need not get PRR neighbors correct, only leaf set
 - If you believe the Chord work, this isn't too hard to do
- Leaf set also gives implementation of $owner^{-1}(n)$
 - All identifiers half-way between n and its predecessor to half-way between n and its successor
- Can store k predecessors and successors
 - Gives further robustness as in Chord

18


Joining a Pastry Network

- Must know of a "gateway" node, g
- Pick new node's identifier, n , U.A.R. from I
- Ask g to find the $m = owner(n)$
 - And ask that it record the path that it takes to do so
- Ask m for its leaf set
- Contact m 's leaf set and announce n 's presence
 - These nodes add n to their leaf sets and vice versa
- Build routing table
 - Get level i of routing table from node i in the join path
 - Use those nodes to make level i of our routing table

19

Pastry Join Example

- Node 3701 wants to join
 - Has 1000 as gateway
- Join path is 1000→3800
 - 3800 is the owner
- 3701 ties itself into leaf set
- 3701 builds routing table
 - L0 neighbors from 1000
 - 1000, 2000, and 3800
 - L1 neighbors from 3800
 - 3600
- Existing nodes on join path consider 3701 as a neighbor



20

Pastry Join Notes

- Join is not “perfect”
 - A node whose routing table needs new node should learn about it
 - Necessary to prove $O(\log |N|)$ routing hops
 - Also not guaranteed to find close neighbors
- Can use *routing table maintenance* to fix both
 - Periodically ask neighbors for their neighbors
 - Use to fix routing table holes; replace existing distant neighbors
- Philosophically very similar to Chord
 - Start with minimum needed for correctness (leaf set)
 - Patch up performance later (routing table)

21

Pastry Join Optimization

- Best if gateway node is “close” to joining node
 - Gateway joined earlier, should have close neighbors
 - Recursively, gateway’s neighbors’ neighbors are close
 - Join path intuitively provides good initial routing table
 - Less need to fix up with routing table maintenance
- Pastry’s optimized join algorithm
 - Before joining, find a good gateway, then join normally
- To find a good gateway, refine set of candidates
 - Start with original gateway’s leaf set
 - Keep only a closest few, then add their neighbors
 - Repeat (more or less--see paper for details)

22

Leaving a Pastry Network

- Do not distinguish between leaving and crashing
 - A good design decision, IMHO
- Remaining nodes notice leaving node n down
 - Stops responding to keep-alive pings
- Fix leaf sets immediately
 - Easy if 2+ predecessors and successors known
- Fix routing table lazily
 - Wait until needed for a query
 - Or until routing table maintenance
 - Arbitrary decision, IMHO

23

Dealing with Broken Leaf Sets

- What if no nodes left in leaf set due to failures?
- Can use routing table to recover (MCR93)
 - Choose closest nodes in routing table to own identifier
 - Ask them for their leaf sets
 - Choose closest of those, recurse
- Allows use of smaller leaf sets

24

Lecture Overview

- Introduction
- PRR Trees
 - Overview
 - Locality Properties
- Pastry
 - Routing in Pastry
 - Joining a Pastry network
 - Leaving a Pastry network
- Tapestry
 - Routing in Tapestry
 - Object location in Tapestry
- Multicast in PRR Trees
- Conclusions

25

Tapestry Routing

- Only different from Pastry when no exact match
 - Instead of using next numerically closer node, use node with next higher digit at each hop
- Example:
 - Given 3 node network: nodes 0700, 0F00, and FFFF
 - Who owns identifier 0000?
- In Pastry, FFFF does (numerically closest)
- In Tapestry, 0700 does
 - From FFFF to 0700 or 0F00 (doesn't matter)
 - From 0F00 to 0700 (7 is next highest digit after 0)
 - From 0700 to itself (no node with digit between 0 and 7)

26

Notes on Tapestry Routing

- Mostly same locality properties as PRR and Pastry
- But compared to Pastry, very fragile
- Consider previous example: 0700, 0F00, FFFF
 - What if 0F00 doesn't know about 0700?
 - 0F00 will think it is the owner of 0000
 - 0700 will still think it is the owner
 - Mapping won't be deterministic throughout network
- Tapestry join algorithm guarantees won't happen
 - All routing table holes than can be filled will be
 - Provably correct, but tricky to implement
 - Leaf set links are bidirectional, easier to keep consistent

27

Object Location in Tapestry

- Pastry was originally just a DHT
 - Support for multicast added later (RKC+01)
- PRR and Tapestry are DOLRs
 - Distributed Object Location and Routing
- Service model
 - `publish (name)`
 - `route_to_object (name, message)`
- Like Napster, Gnutella, and DNS
 - Service does not store data, only pointers to it
 - Manages a mapping of names to hosts

28

A Simple DOLR Implementation

- Can implement a DOLR on owner service model

```
publish (name) {
    n = owner (name)
    n.add_mapping (name, my_addr, my_port)
}

route_to_object (name, message) {
    n = owner (name)
    m = n.get_mapping (name)
    m.send_msg (message)
}
```

29

Problems with Simple DOLR Impl.

- No locality
 - Even if object stored nearby, owner might be far away
 - Bad for performance
 - Bad for availability (owner might be behind partition)
- No redundancy
 - Easy to fix if underlying network has leaf/successor sets
 - Just store pointers on owner's whole leaf set
 - If owner fails, replacement already has pointers
 - But Tapestry doesn't have leaf/successor sets

30

Tapestry DOLR Implementation


- Insight: leave "bread crumbs" along publish path
 - Not just at owner

```
publish (name) {
    foreach n in path_to_owner (name)
        n.add_mapping (name, my_addr, my_port)
}

route_to_object (name, message) {
    foreach n in path_to_owner (name)
        if ((m = n.get_mapping (name)) != null) {
            m.send_msg (message); break;
        }
}
```

31

Tapestry DOLR Impl.: Experiments




32

Tapestry DOLR Impl. Notes

- Bread crumbs called “object pointers”
- PRR show that overlay path from query source to object is no worse than a constant longer than underlying network path
 - Just like in routing in a PRR tree
- True for two reasons:
 1. Hops early in path are short (in network latency)
 2. Paths converge early
- Path convergence is a little subtle
 - Two nearby nodes often have same early hops
 - Because next hop based on destination, not source
 - And because neighbor choice weighted on latency

33

Path Convergence Examples



34

Lecture Overview

- Introduction
- PRR Trees
 - Overview
 - Locality Properties
- Pastry
 - Routing in Pastry
 - Joining a Pastry network
 - Leaving a Pastry network
- Tapestry
 - Routing in Tapestry
 - Object location in Tapestry
- Multicast in PRR Trees
- Conclusions

35

Multicast in PRR Trees

- PRR Tree gives efficient paths between all nodes
 - Uses application nodes as routers
- Since control routers, can implement multicast
 - Seems to have been thought of simultaneously by:
 - Pastry group, with SCRIBE protocol (RKC+01)
 - Tapestry group, with Bayeux protocol (ZZJ+01)
- I'll talk about SCRIBE


36

SCRIBE Protocol

- Like Tapestry object location, use bread crumbs
- To subscribe to multicast group i
 - Walk up tree towards *owner* (i), leaving bread crumbs
 - If find pre-existing crumb, leave one more and stop
- To send a message to group i
 - Send message to *owner* (i)
 - Owner sends message to each bread crumb it has
 - Its children who are subscribers or their parents
 - Each child recursively does the same

37

SCRIBE Example



38

SCRIBE Notes

- Multicast evaluation points
 - Stretch, also called Relative Delay Penalty (RDP)
 - Network Stress
- Scribe has constant RDP
 - Assuming Pastry is a good PRR tree
- Scribe has low network stress
 - Harder to see, but due to choosing neighbors by latency
 - Demonstrated in simulations (CJK03)

39

Conclusions

- PRR trees are a powerful P2P primitive
 - Can be used as a DHT
 - Path to owner has low RDP
 - Chord can be "hacked" to do the same
 - Can be used as a DOLR
 - Finds close objects when available
 - No clear way to get this from Chord
 - Can be used for application-level multicast
 - Also no clear way to get this from Chord
- More work to support dynamic membership
 - Pastry uses leaf sets like Chord
 - Tapestry has own algorithm

40

For more information

- Pastry
<http://research.microsoft.com/~antr/pastry/pubs.htm>
- Tapestry
<http://oceanstore.cs.berkeley.edu/publications/index.html>

41